

Written Interview for Python Engineer - Data Center Hardware Integration (Greater Boston Area)

Name Redacted

October 10, 2023

1 Career Development

1.1 How would you describe your level of experience managing data servers and network hardware?

I have an intermediate level of experience managing data servers and network hardware. I have performed most common tasks related to configuring and maintaining data servers and network hardware at least once. In the past, I was a Cisco Certified Entry Network Technician certification and am familiar with configuring hardware running Cisco IOS.

In both my graduate and undergraduate research labs, I received training in configuring and installing rack-mounted Linux servers. This training was used for installing a handful of hardware, mostly GPU servers. This involved everything from physical installations to account management to software installation.

1.2 How would you describe your level of experience as a software engineer?

I believe that my experience has me on track to being a senior engineer in the next 12-18 months. My prior experiences as a researcher, machine learning engineer, and consultant have honed my abilities as a software engineer who had to ship code to production and could be easily maintained by other engineers. I am comfortable writing code in multiple languages, with my strongest being Python. Here is an abbreviated list of professional software projects I have worked on in the past:

- Topic classification microservice for a microblog-based social media platform: Python
- Optimized post-processing pipeline for real-time object detection in video: Python, C++17, CUDA
- A cycle-approximate simulation of a RISC-V CPU down to the basic block level: Golang

1.3 What is your proudest success as an engineer?

To this day, my proudest success is the Discord bot which plays background music for my friends and while we hang out in voice chat. While it is not my most technically advanced creation, it is something I made which improves the lives of people that are important to me. I often get requests for feature additions, or bug reports from my friends even when I'm not online, which tells me that they use it even when I'm not around.

1.4 Outline the role of an engineering manager in shaping a high functioning team.

An engineering manager shapes a high functioning team in the following ways:

- **Hiring and Firing:** The Engineering Manager is the most aware of what skills would be needed to complement the existing team, and thus they should have a large role in selecting new hires for their team.

- **Managing Upwards:** The Engineering Manager should have the most time to meet with leadership, as implementation work has been delegated to the engineers that they manage. This time should be used to ensure that their engineers have all the resources (e.g., time, information, infrastructure) they need to do well, and that expectations are managed accordingly
- **Mentorship:** The Engineering Manager should serve as a "force multiplier" for the engineers they manage, by allocating work in such a way that the engineers can grow technically and professionally.

2 Experience

2.1 Describe your level of experience in Python, and how you have attained it.

2.2 Describe a case where it was very difficult to test code you were writing, but you found a reliable way to do it.

While working as a Machine Learning Engineer, it was common to encounter code which compiled and ran without software errors, but didn't perform as expected on certain edge cases. One of the first tasks I set about doing as a new hire was to develop a special test suite that would allow for us to measure whether changes to our image processing code affected the results of the object classification software (In terms of AUC ROC[3]). This way, every time a change was made to the image processing code for the classifiers, we could observe any potentialy degradation in our image detection capabilities and prevent it from being deployed.

2.3 When did you start working with Linux? Describe your level of experience as a user & developer on Linux.

I started using Linux on a regular basis 8 years ago for my computer science coursework and have been using it ever since. I've used a variety of Linux distros for software development, hosting, to media playback and gaming. I used a Linux distro to write this document. Ubuntu was the Linux distribution of choice for the computing department at both my undergraduate and graduate institutions.

2.4 Describe your experience with networking, storage technologies and cloud infrastructure

I acquired some professional experience with cloud infrastructure by nature of the small engineering team at my previous employer. This involved the following:

- Writing Dockerfiles which assembled a custom SDK for `x64` and `aarch64` architectures
- Managing Artifacts in Google Cloud Platform Artifact Registry
- Writing custom CI workflows for GitHub CI and Jenkins

2.5 Describe your experience of large-scale physical server installations, including any provisioning, automation and service orchestration work.

2.6 Tell us what you believe teams should consider when they build, test, run and deliver software.

2.7 How do you think about quality?

Quality is what the you remember after using a product or service. It's what provides the sensation of "wow, this really just works" when using a piece of software. Humans are naturally highly perceptive creatures, and thus high quality products and services require lots of attention to the finest details.

2.8 What would you like to achieve in career development and skills development?

3 Education

3.1 At high school, how did you fare in mathematics and physical sciences? Which were your strongest subjects in the hard sciences, and how did you rank in your class?

I was very strong in math and physical sciences in high school. I competed at the state level in both domains, and was ranked 18th out of a class of 496 (Top 5%).

3.2 At high school, what leadership roles did you take on?

Here is a list of relevant leadership roles from my high school years:

- I was the lead editor for all Math and computing content at my highschool's educational newsletter.
- I tutored math and Spanish for 4 hours a week as a volunteer.
- I represented my high school at the University of South Carolina Summer at Moore business competition.

3.3 What course and university did you choose, and why?

I attended Clemson University and acquired a Bachelors of Science in Computer Science, because understanding how computers work has been my passion since a very young age. I was awarded the Palmetto Fellows Scholarship (along with other awards), and thus was given the opportunity to attend at very low cost, cementing my choice of university.

3.4 How did you rank competitively in university? Which were your strongest courses, and which did you enjoy the most?

I was mostly an A/B student in my undergraduate studies, likely in the upper 30% of the class. Here is a sample of courses where I recieved the highest grade in my undergraduate:

- Network Programming
- Design and Analysis of Algorithms
- Calculus of Several Variables

I really enjoyed courses that allowed me to take skills I had developed earlier on and build projects, here are some examples:

- 2-D Game Engine Construction
- Applied Data Science
- Robotics

My academic showing during my Master of Science in Computer Engineering was much stronger, with a GPA of 3.8/4.0.

3.5 At high school and university, describe your achievements that were considered exceptional by colleagues and staff.

- **Clemson School of Computing Best Research** I built monitoring infrastructure for memory accesses in Nvidia's Unified Virtual Memory system, and eventually presented the research at Super Computing 2019.
- **1st Place South Carolina Junior Academy of Sciences, Math and Computing** I studied reward allocation behavior in Bitcoin mining pools from 2013-2014 to determine if certain policies were more beneficial to small, individual mining participants such as myself.
- **LQTLDD3** I wrote a map-parsing library for a robotics project that allowed for the path-finding algorithm to determine available adjacent space in constant time, and then made it open source. My implementation is based a paper by Kunio Aizawa and Shojiro Tanaka[1] and the repository is hosted on GitHub[4].

4 Context

4.1 How are you involved in open source software? Describe any significant contributions to open source (with links where possible)

As a professional developer, I am a contributor to Akita[5], which is an MIT-licensed CPU/GPU simulation framework for building and testing microarchitectures against different workloads. I became a contributor for Akita through my graduate research, where I used Akita to build a RISC-V emulator.

As a user, I see it as my duty to file issues and work on improving the product when possible, however previous employers chose not allocate time towards merging our code upstream in cases where we did work with open source software.

4.2 What do you think are the key ingredients of a successful open source project?

- **Effective leadership**, who is capable of communicating priorities based on a strong vision, and running the project well at it's current stage.
- **Passionate builders**, who are technically capable to convert priorities into features. Often in the early stages, the leadership is also building.
- **Transparency**. Open source software wins because it is extensible and transparent. It is easiest to build on software when the person building on top of the software has all of the control they need to get the job done. There was a great essay on this recently by the Co-founder of Lago[2].

4.3 Why do you most want to work for Canonical?

I firmly believe in Canonical's mission to make the frontier of technology an accessible resource to all. I originally applied to grad school so that I could help make advanced machine learning methods more accessible. Building infrastructure with Canonical is a solid step towards making my mission a reality.

4.4 Which other companies are building the sort of products you would like to work on?

Here's what I seek out in terms of projects:

- Technical Challenge to improve / maintain
- Provides utility to people

- Generally related to a subject I'm experienced with (e.g., Hardware/Software, Audio/Video, Machine Learning)

Here are some examples of companies that make products which meet those requirements:

- Apple (WebKit, Clang/LLVM)
- Microsoft (GitHub, Visual Studio Code)
- HuggingFace (HuggingFace Hub)

4.5 What do you think Canonical needs to improve in its engineering and products?

4.6 Who do you think are key competitors to Canonical? How do you think Canonical should plan to win that race?

The key competitors to Canonical are RedHat/IBM, SUSE, and to a lesser extent, HashiCorp. These companies compete with Canonical by providing software which attempts to simplify and streamline managing enterprise software + infrastructure. I think beating out this portion of the competition comes down to providing a superior UX on the handful of workflows that matter the most to a target userbase.

References

- [1] Kunio Aizawa and Shojiro Tanaka. A constant-time algorithm for finding neighbors in quadrees. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 31(7):1178–1183, 2009.
- [2] Anh-Tho Chong. Open source does not win by being cheaper. Retrieved from URL: <https://github.com/getlago/lago/wiki/Open-Source-does-not-win-by-being-cheaper>, 2023.
- [3] Google Machine Learning Education. Classification: Roc curve and auc. Foundational Courses: Crash Course, 2022. Retrieved from URL: <https://developers.google.com/machine-learning/crash-course/classification/roc-and-auc>.
- [4] Derek Rodriguez. Lqtld3. Retrieved from URL: <https://github.com/dwrodri/LQTLD3>, 2017.
- [5] Yifan Sun. Akita (github repo). Retrieved from URL: <https://github.com/sarchlab/akita>, 2018.